

Generative Adversarial Networks

Egor Zakharov

Engineer at Samsung AI Center


PhD student at Skolkovo Institute of Science and Technology




Outline

- Part I
 - Intuition behind GANs
 - Implicit vs explicit generative models
 - JS and f-divergences
 - Wasserstein GAN
- Part II
 - Spectral normalization
 - Moving averaging, weight averaging
 - Quality measurements
 - Applications


Classification




Classification




Classification



Classification




Classification




Some intuition

- So far we had fixed $p(x)$, $q(x)$ and only trained classifier
- How do we use classifier's output to move $q(x)$ towards $p(x)$?

Using classifier's feedback



Using classifier's feedback



Parametrization

1. What do we need to learn a classifier?

Only samples from $p(x)$ and $q(x)$!

2. How do we parametrize model distribution $q(x)$?

Parametrize density function

e. g. $q_{\theta}(x) = \mathcal{N}(x; \theta, I)$

- We should be able to sample from q_{θ}
- Have access to density at any point.

Define implicitly

$$z \sim \mathcal{N}(0, I), \quad G_{\theta}(z) \sim q_{\theta}(x)$$

- Sampling is always easy
- Hard to evaluate point density $q_{\theta}(x)$

In case of images


Noise $\sim N(0,1)$



Generative
Model



Simulation



GAN Game

- Classifier

$$f_{\phi}(\mathbf{x}) = p_{\phi}(y = 1 | \mathbf{x})$$

- Classification loss

$$\mathcal{L}(\phi, \theta) = \mathbb{E}_{p(\mathbf{x})}[-\log f_{\phi}(\mathbf{x})] + \mathbb{E}_{p(\mathbf{z})}[-\log(1 - f_{\phi}(G_{\theta}(\mathbf{z})))]$$

Algorithm

1. **Update classifier**

$$\phi^* = \arg \min_{\phi} \mathcal{L}(\phi, \theta)$$

2. **Update generator**

$$\theta^{new} = \theta^{old} + \frac{\partial \mathcal{L}(\phi^*, \theta^{old})}{\partial \theta}$$

3. **Repeat**

In general

General learning scheme:

1. Update guide

- $\frac{p(\mathbf{x})}{p(\mathbf{x}) + q(\mathbf{x})}$
- $\frac{p(\mathbf{x})}{q(\mathbf{x})}$
- $p(\mathbf{x}) - q(\mathbf{x})$
- $\mathcal{D}(p||q)$


2. Use guide to **update generator**

- Move $q(\mathbf{x})$ closer to $p(\mathbf{x})$

3. Repeat

Prescribed vs implicit models

Prescribed (think of VAE)

- $p(z)$
 - $q(x)$
 - $p(x|z)$
 - $q(z|x)$
 - $q(x, z)$
- 
- Evaluate and sample**

Implicit (think of GAN)

- Evaluate and sample from $p(z)$
- Sample from $p(x), q(x)$
- Approximate $q(x)$ using samples
- Approximate $q(z|x)$

GAN Game

- Classification loss

$$\mathcal{L}(\phi, \theta) = \mathbb{E}_{p(\mathbf{x})}[-\log f_{\phi}(\mathbf{x})] + \mathbb{E}_{p(\mathbf{z})}[-\log(1 - f_{\phi}(G_{\theta}(\mathbf{z})))]$$

Algorithm

1. **Update classifier**

$$\min_{\phi} \mathcal{L}(\phi, \theta)$$

2. **Update generator**

$$\theta^{new} = \theta^{old} + \frac{\partial \mathcal{L}(\phi^*, \theta^{old})}{\partial \theta}$$

3. **Repeat**

GAN Game

- Classification loss

$$\mathcal{L}(\phi, \theta) = \mathbb{E}_{p(\mathbf{x})}[-\log f_{\phi}(\mathbf{x})] + \mathbb{E}_{p(\mathbf{z})}[-\log(1 - f_{\phi}(G_{\theta}(\mathbf{z})))]$$

Algorithm

1. **Update classifier**

$$\max_{\phi} -\mathcal{L}(\phi, \theta)$$

2. **Update generator**

$$\theta^{new} = \theta^{old} + \frac{\partial \mathcal{L}(\phi^*, \theta^{old})}{\partial \theta}$$

3. **Repeat**

GAN Game

- Classification loss

$$\mathcal{L}(\phi, \theta) = \mathbb{E}_{p(\mathbf{x})}[-\log f_{\phi}(\mathbf{x})] + \mathbb{E}_{p(\mathbf{z})}[-\log(1 - f_{\phi}(G_{\theta}(\mathbf{z})))]$$

Algorithm

1. **Update classifier**

$$\max_{\phi} -\mathcal{L}(\phi, \theta) = -\log(4) + 2\mathcal{D}_{JS}(p||q_{\theta})$$

2. **Update generator**

$$\theta^{new} = \theta^{old} + \frac{\partial \mathcal{L}(\phi^*, \theta^{old})}{\partial \theta}$$

3. **Repeat**

GAN Game

- Classification loss

$$\mathcal{L}(\phi, \theta) = \mathbb{E}_{p(\mathbf{x})}[-\log f_{\phi}(\mathbf{x})] + \mathbb{E}_{p(\mathbf{z})}[-\log(1 - f_{\phi}(G_{\theta}(\mathbf{z})))]$$

Algorithm

1. **Update classifier**

Estimate (kind of) distance between $p(\mathbf{x})$ and $q(\mathbf{x})$

$$\max_{\phi} -\mathcal{L}(\phi, \theta) = -\log(4) + 2\mathcal{D}_{JS}(p||q_{\theta})$$

2. **Update generator**

$$\theta^{new} = \theta^{old} + \frac{\partial \mathcal{L}(\phi^*, \theta^{old})}{\partial \theta}$$

3. **Repeat**

Minimize the distance

(GAN) Game

- (Classification) loss

$$\mathcal{L}(\phi, \theta) = ?$$

Algorithm

1. **Update classifier**

Estimate (kind of) distance between $p(\mathbf{x})$ and $q(\mathbf{x})$

$$\max_{\phi} -\mathcal{L}(\phi, \theta) = \mathcal{D}(p \| q_{\theta})$$

2. **Update generator**

$$\theta^{new} = \theta^{old} + \frac{\partial \mathcal{L}(\phi^*, \theta^{old})}{\partial \theta}$$

3. **Repeat**

Minimize the distance

f-Divergence

- For distributions P and Q **f-divergence** is defined as:

$$D_f(P \parallel Q) = \int_{\mathcal{X}} f\left(\frac{p(x)}{q(x)}\right) q(x) \, dx,$$

- KL-divergence:** $f(t) = t \log(t)$

$$D_f(P \parallel Q) = KL(P \parallel Q)$$

- Reversed KL-divergence:** $f(t) = -\log(t)$

$$D_f(P \parallel Q) = KL(Q \parallel P)$$

- Total variation:** $f(t) = \frac{1}{2}|t - 1|$

$$D_f(P \parallel Q) = \frac{1}{2} \int_{\mathcal{X}} |p(x) - q(x)| \, dx$$

Optimal transport




- Define a **cost** of transporting from x to y as $c(x, y)$
 - e.g. $c(x, y) = ||x - y||$
- **Optimal transport** cost is then defined as:

$$T(P, Q) = \inf_{\Gamma \in \mathcal{P}(x \sim P, y \sim Q)} \mathbb{E}_{(x, y) \sim \Gamma} [c(x, y)]$$

- where $\mathcal{P}(x \sim P, y \sim Q)$ is a set of all joint distributions of (x, y) with marginals P and Q respectively.

Optimal transport: example



Optimal transport dual

- **Primal:**

$$T(P, Q) = \inf_{\Gamma \in \mathcal{P}(x \sim P, y \sim Q)} \mathbb{E}_{(x,y) \sim \Gamma} [c(x, y)]$$

- **Dual (Wasserstein-1 metric):**

$$T(P, Q) = W_1(P, Q) = \sup_{\|f\|_L \leq 1} \mathbb{E}_{x \sim P} f(x) - \mathbb{E}_{x \sim Q} f(x)$$


How to satisfy lipschitz continuity condition?



Two options:

- Weight clipping (used in original WGAN paper)
- Gradient penalty $\lambda \mathbb{E}_{\tilde{x} \sim \mathbb{P}(\tilde{x})} (\|\nabla_{\tilde{x}} D(\tilde{x})\|_2 - 1)^2$

Optimal transport vs f-Divergence



- $W_1(P, Q) = \theta$

- $JS(P\|Q) = \begin{cases} \log(2), & \theta \neq 0 \\ 0, & \theta = 0 \end{cases}$

- $KL(P\|Q) = \begin{cases} \infty, & \theta \neq 0 \\ 0, & \theta = 0 \end{cases}$

Outline

- Part I
 - Intuition behind GANs
 - Implicit vs explicit generative models
 - JS and f-divergences
 - Wasserstein GAN

- Part II
 - Spectral normalization
 - Moving averaging, weight averaging
 - Quality measurements
 - Applications

Spectral Normalization

- On practice, gradient penalty $\lambda \mathbb{E}_{\tilde{x} \sim \mathbb{P}(\tilde{x})} (\|\nabla_{\tilde{x}} D(\tilde{x})\|_2 - 1)^2$
 - is a restrictive constraint
 - introduces significant bias
- Is there an easier way to introduce Lipchitz-1 continuity?
 - Reminder of how the convolutions work:

3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Convolutions as linear operators

$$\begin{array}{ccc} \begin{pmatrix} x_1 & x_2 & x_3 \\ x_4 & x_5 & x_6 \\ x_7 & x_8 & x_9 \end{pmatrix} & \begin{pmatrix} k_1 & k_2 \\ k_3 & k_4 \end{pmatrix} & x \circ \omega = ? \\ x & \omega & \end{array}$$

$$x \circ \omega = \begin{pmatrix} x_1 k_1 + x_2 k_2 + x_4 k_3 + x_5 k_4 & x_2 k_1 + x_3 k_2 + x_5 k_3 + x_6 k_4 \\ x_1 k_4 + x_5 k_2 + x_7 k_3 + x_8 k_4 & x_5 k_1 + x_6 k_2 + x_8 k_3 + x_9 k_4 \end{pmatrix}$$

Convolutions as linear operators

$$\begin{pmatrix} k_1 & k_2 & 0 & k_3 & k_4 & 0 & 0 & 0 & 0 \\ 0 & k_1 & k_2 & 0 & k_3 & k_4 & 0 & 0 & 0 \\ 0 & 0 & 0 & k_1 & k_2 & 0 & k_3 & k_4 & 0 \\ 0 & 0 & 0 & 0 & k_1 & k_2 & 0 & k_3 & k_4 \end{pmatrix} \cdot \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{pmatrix} = \begin{pmatrix} x_1 k_1 + x_2 k_2 + x_4 k_3 + x_5 k_4 \\ x_2 k_1 + x_3 k_2 + x_5 k_3 + x_6 k_4 \\ x_1 k_4 + x_5 k_2 + x_7 k_3 + x_8 k_4 \\ x_5 k_1 + x_6 k_2 + x_8 k_3 + x_9 k_4 \end{pmatrix}$$

A_ω
 \tilde{x}

Lipchitz continuity of a linear operator

$$\|f(x_1) - f(x_2)\|_2 \leq L \|x_1 - x_2\|_2$$

$$\|Ax_1 - Ax_2\|_2 \leq L \|x_1 - x_2\|_2$$

$$\frac{\|Ax\|_2}{\|x\|_2} \leq L$$

$$\sigma_{\max} = \sup_x \frac{\|Ax\|_2}{\|x\|_2} \leq L$$

Spectral norm

Singular values and where to find them

- All singular values can be calculated via the singular value decomposition (SVD):

$$\begin{array}{ccc} \text{"left" singular vectors} & \rightarrow & \\ & A = U \Sigma V^T & \leftarrow \\ & & \text{"right" singular vectors} \end{array}$$

- Given a "left" and "right" singular vectors u and v , the corresponding singular value equals to:

$$\sigma = u^T A v$$

- Knowing that, there is a simpler way to get a maximum singular value

Power iteration

- For simplicity, assume that A is square and diagonalizable (instead of an SVD we can do a simpler EVD: $A = Q\Lambda Q^{-1}$)

- Let x_1 be initialized at random and have a unit norm
- Run the following iteration:

1. $x_{i+1} = Ax_i$

2. $x_{i+1} = \frac{x_{i+1}}{\|x_{i+1}\|_2}$

- In the end, we obtain an eigenvector q , corresponding to a maximum eigenvalue λ_{\max}

$$\lambda_{\max} = q^T A q$$

Theory vs practice

- In practice, instead of A_ω , a reshaped version W of the weights ω is used

$$\begin{array}{ccc} & \uparrow & \uparrow \\ C_{\text{out}}H_{\text{out}}W_{\text{out}} \times C_{\text{in}}H_{\text{in}}W_{\text{in}} & & C_{\text{out}} \times C_{\text{in}}K_hK_w \end{array}$$

- Can be proven that the resulting $\tilde{\sigma}_{\text{max}}$ is a lower bound to the true σ_{max}
- This method consistently outperforms “true” spectral normalization
- Takeaway: explanation in the original paper is at least incomplete, further research is needed


Game theory view on GANs

$$\min_{\theta} \max_{\psi} \underbrace{\mathbb{E}_{x \sim p(x)} [\log D_{\psi}(x)] + \mathbb{E}_{z \sim p(z)} [\log(1 - D_{\psi}(G_{\theta}(z)))]}_{\mathcal{L}(\theta, \phi)}$$

In theory:

1. $\psi^* = \arg \max_{\psi} \mathcal{L}(\theta^{\text{old}}, \psi)$
2. $\theta^{\text{new}} = \theta^{\text{old}} - \alpha \nabla_{\theta} \mathcal{L}(\theta^{\text{old}}, \psi^*)$
3. $\theta^{\text{old}} = \theta^{\text{new}}$
4. Go to step 1

$\mathcal{D}(P||Q)$



In practice:

1. $\psi^{\text{new}} = \psi^{\text{old}} + \alpha \nabla_{\psi} \mathcal{L}(\theta^{\text{old}}, \psi^{\text{old}})$
2. $\theta^{\text{new}} = \theta^{\text{old}} - \alpha \nabla_{\theta} \mathcal{L}(\theta^{\text{old}}, \psi^{\text{new}})$
3. $\theta^{\text{old}} = \theta^{\text{new}}, \psi^{\text{old}} = \psi^{\text{new}}$
4. Go to step 1

Example

Is the saddle point reachable?


$$\min_{\theta} \max_{\psi} \mathbb{E}_{x \sim p(x)} f(D_{\psi}(x)) + \mathbb{E}_{z \sim p(z)} f(1 - D_{\psi}(G_{\theta}(z)))$$

$$\min_{\theta} \max_{\psi} \mathbb{E}_{x \sim \delta_0} f(\psi \cdot x) + \mathbb{E}_{\hat{x} \sim \delta_{\theta}} f(\psi \cdot \hat{x})$$


$$\min_{\theta} \max_{\psi} f(0) + f(\psi\theta)$$

$$\min_{\theta} \max_{\psi} f(\psi\theta)$$

Is the saddle point reachable?



Non-saturating GAN




WGAN-GP

Exponential moving average

$$\hat{\theta}_{i+1} = (1 - \alpha)\hat{\theta}_i + \alpha\theta_{i+1}$$


- Reduces the radius of the orbit
- Works “out of the box” (momentum is adjusted according to the planned training time)



Weight averaging

$$\hat{\theta} = \frac{1}{N - n} \sum_{i=n}^N \theta_i$$

- Averaging the weights over the orbit gets you directly to the saddle point
- Problem: when to start averaging?




Quality measurements

- Assume that you have trained a GAN, how to tell if it's good?
- All automated metrics suffer from being “bad”, i.e. good models may get low scores
- Best metrics are usually hand-crafted and problem-specific
 - Example: perceptual image quality may be measured via a user study
- Let's discuss some metrics used for the images

Inception Score (IS)

- Evaluates “objectiveness” and class distribution
 - Assumption: each generated image has a single object in it
 - The overall distribution of the classes has to be uniform



Inception Score (IS)



low entropy $p(y|x)$

high entropy $p(y) = \sum_x p(y|x)p(x)$

$$\begin{aligned}
 \log \text{IS} &= \mathbb{E}_x \text{KL}(p(y|x) || p(y)) = \sum_x p(x) \sum_y p(y|x) \log \frac{p(y)}{p(y|x)} \\
 &= \underbrace{\sum_y \log p(y) \sum_x p(y|x)p(x)}_{H(y)} - \underbrace{\sum_x \sum_y p(y|x)p(x) \log p(y|x)}_{H(y|x)}
 \end{aligned}$$

Frechet Inception Distance (FID)

- Measures the quality of the generated images
- tl;dr; compare the activations of a pre-trained convnet between real and generated datasets
- We compare only means and pairwise correlations

$$\text{FID} = ||\mu_r - \mu_g||^2 + \text{Tr}(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{1/2})$$

Why you may need GANs



2014



2015



2016



2017



2018

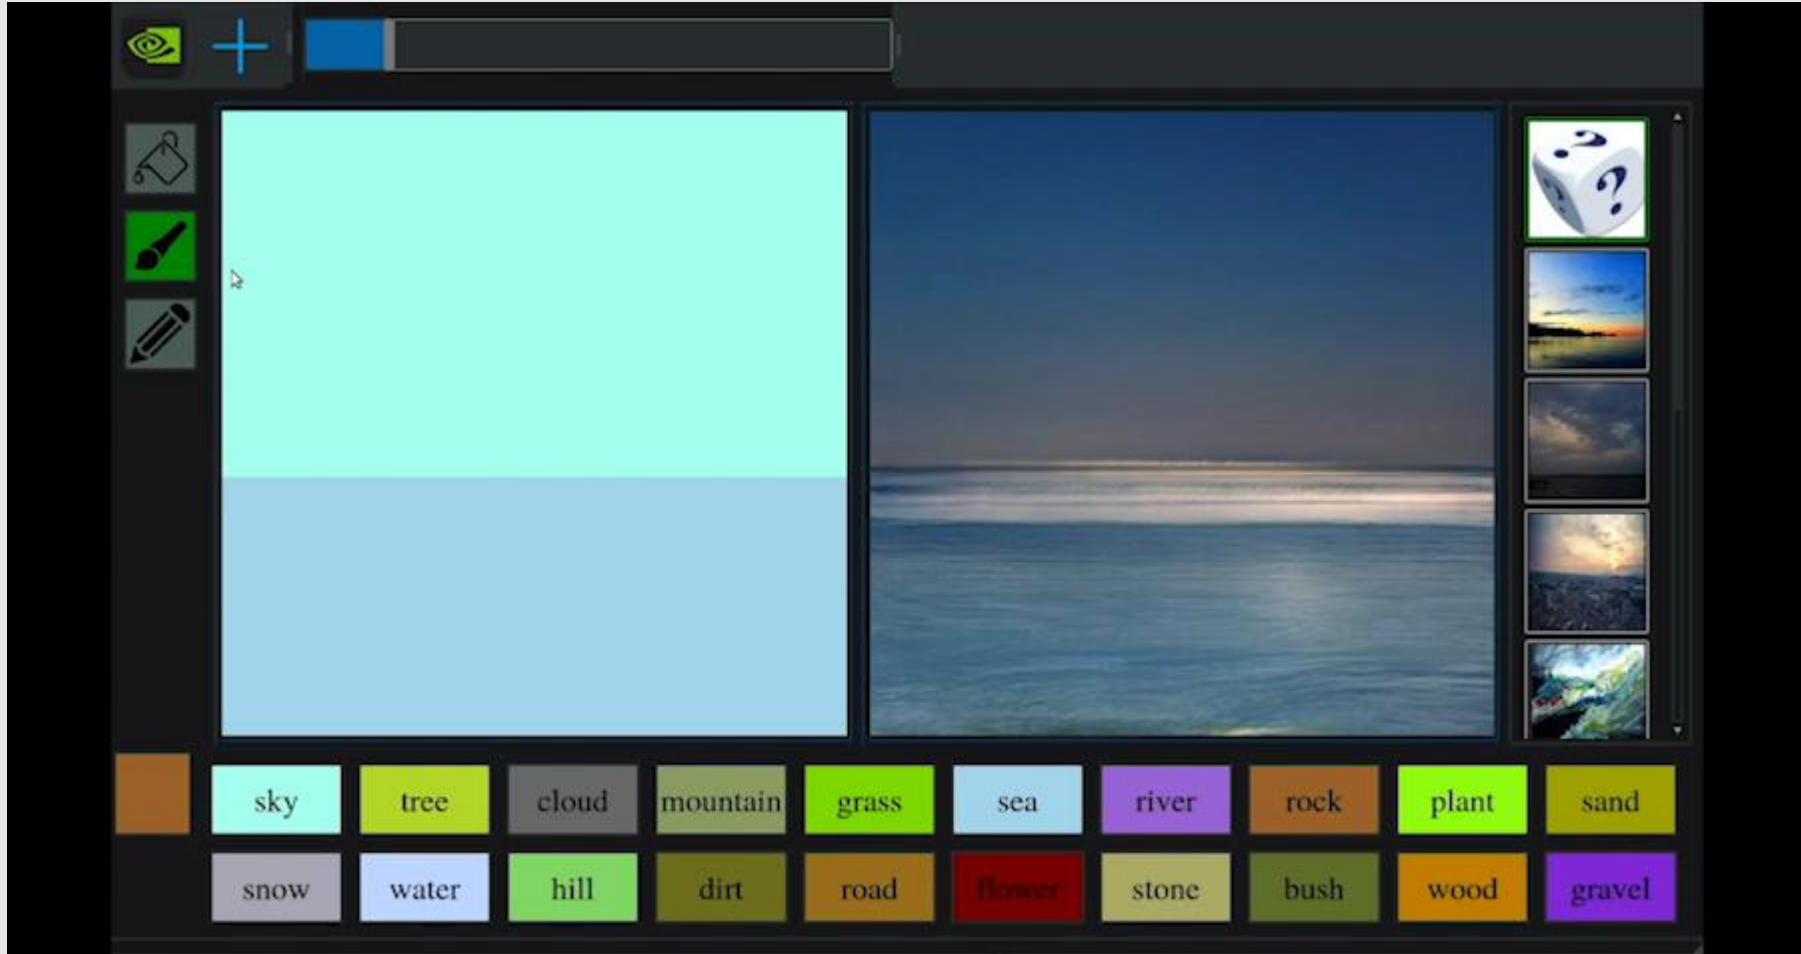
Images credit: Ian Goodfellow, twitter.com/goodfellow_ian

Why you may need GANs



Image credit: Peter Baylies, twitter.com/pbdaylies

Why you may need GANs



Park et al., "Semantic Image Synthesis with Spatially-Adaptive Normalization"

Takeaway message

- GANs are powerful implicit generative models that achieve SotA “quality” of samples
- Recent works show that the improvement in quality of the “outputs” corresponds to improvements in the trained latent space
- Some intuitive explanations of why the GANs work may be incomplete
- Application of GANs and related techniques go far beyond the problems of sampling